



UNIVERSIDADE DE COIMBRA

Faculdade de Ciências e Tecnologias,
Departamento de Engenharia Informática,

Mestrado em Engenharia e Ciência de Dados

Advanced Infrastructures for Data Science

Project #2

João Pino, 2020210945

Luís Seco, 2020211927

Assignment 1 - Kafka

Setup Zookeper and Kafka Broker

- zkServer.sh start
- cd
- cd kafka-3.9.0-src/
- bin/kafka-server-start.sh config/server.properties

Exercise 1

Terminal 1

- source venv/bin/activate
- python3 consumer.py

```
(venv) joaopino@din-org39 task1-topic %
python3 consumer.py
Listening to task1-topic...
Received: {'sensor_id': 4, 'temperature': 28.07}
Received: {'sensor_id': 8, 'temperature': 27.93}
Received: {'sensor_id': 4, 'temperature': 31.77}
Received: {'sensor_id': 9, 'temperature': 34.05}
```

Terminal 2

- source venv/bin/activate
- python3 producer.py

```
(venv) joaopino@din-org39 task1-topic %
python3 producer.py
Sent: {'sensor_id': 4, 'temperature': 28.07}
Sent: {'sensor_id': 8, 'temperature': 27.93}
Sent: {'sensor_id': 4, 'temperature': 31.77}
Sent: {'sensor_id': 9, 'temperature': 34.05}
```

Exercise 2

- Create a topic-2 with 3 partitions:

```
bin/kafka-topics.sh --create \
--topic task2-topic \
--bootstrap-server localhost:9092 \
--partitions 3 \
--replication-factor
```

```
(venv) joaopino@din-org39 task2-topic % python3 consumer.py
Listening to task2-topic as part of activity-group...
Received: {'user_id': 'user28', 'activity': 'login'}
Received: {'user_id': 'user0', 'activity': 'add_to_cart'}
Received: {'user_id': 'user4', 'activity': 'add_to_cart'}
Received: {'user_id': 'user9', 'activity': 'login'}
Received: {'user_id': 'user10', 'activity': 'login'}
Received: {'user_id': 'user13', 'activity': 'purchase'}
Received: {'user_id': 'user14', 'activity': 'view_item'}
Received: {'user_id': 'user16', 'activity': 'login'}
Received: {'user_id': 'user23', 'activity': 'logout'}
^CTraceback (most recent call last):
```

```
(venv) joaopino@din-org39 task2-topic % python3 consumer.py
Listening to task2-topic as part of activity-group...
Received: {'user_id': 'user5', 'activity': 'logout'}
Received: {'user_id': 'user17', 'activity': 'view_item'}
Received: {'user_id': 'user20', 'activity': 'view_item'}
```

```
(venv) joaopino@din-org39 task2-topic % python3 producer.py
Sent: {'user_id': 'user0', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user1', 'activity': 'logout'}
Sent: {'user_id': 'user2', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user3', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user4', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user5', 'activity': 'logout'}
Sent: {'user_id': 'user6', 'activity': 'purchase'}
Sent: {'user_id': 'user7', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user8', 'activity': 'login'}
Sent: {'user_id': 'user9', 'activity': 'login'}
Sent: {'user_id': 'user10', 'activity': 'login'}
Sent: {'user_id': 'user11', 'activity': 'login'}
Sent: {'user_id': 'user12', 'activity': 'login'}
Sent: {'user_id': 'user13', 'activity': 'purchase'}
Sent: {'user_id': 'user14', 'activity': 'view_item'}
Sent: {'user_id': 'user15', 'activity': 'logout'}
Sent: {'user_id': 'user16', 'activity': 'login'}
Sent: {'user_id': 'user17', 'activity': 'view_item'}
Sent: {'user_id': 'user18', 'activity': 'logout'}
Sent: {'user_id': 'user19', 'activity': 'logout'}
Sent: {'user_id': 'user20', 'activity': 'view_item'}
Sent: {'user_id': 'user21', 'activity': 'login'}
Sent: {'user_id': 'user22', 'activity': 'view_item'}
Sent: {'user_id': 'user23', 'activity': 'logout'}
Sent: {'user_id': 'user24', 'activity': 'view_item'}
Sent: {'user_id': 'user25', 'activity': 'view_item'}
Sent: {'user_id': 'user26', 'activity': 'add_to_cart'}
Sent: {'user_id': 'user27', 'activity': 'logout'}
Sent: {'user_id': 'user28', 'activity': 'logout'}
Sent: {'user_id': 'user29', 'activity': 'view_item'}
Sent: {'user_id': 'user30', 'activity': 'login'}
^CTraceback (most recent call last):
```

```
(venv) joaopino@din-org39 task2-topic % python3 consumer.py
Listening to task2-topic as part of activity-group...
```

In this project, as we can see, 3 of the consumers are receiving information, one at a time, while 1 of them is not receiving any.

That is due to the topic only having three partitions, so only 3 consumers can actively listen at a time.

We also analysed that in the case of failure in one of the consumers, the fourth one that was actively listening starts receiving information.

```
(venv) joaopino@din-org39 task2-topic % python3 consumer.py
Listening to task2-topic as part of activity-group...
Received: {'user_id': 'user1', 'activity': 'logout'}
Received: {'user_id': 'user2', 'activity': 'add_to_cart'}
Received: {'user_id': 'user3', 'activity': 'add_to_cart'}
Received: {'user_id': 'user6', 'activity': 'purchase'}
Received: {'user_id': 'user7', 'activity': 'add_to_cart'}
Received: {'user_id': 'user8', 'activity': 'login'}
Received: {'user_id': 'user11', 'activity': 'login'}
Received: {'user_id': 'user12', 'activity': 'login'}
Received: {'user_id': 'user15', 'activity': 'logout'}
Received: {'user_id': 'user18', 'activity': 'logout'}
Received: {'user_id': 'user19', 'activity': 'logout'}
Received: {'user_id': 'user21', 'activity': 'login'}
Received: {'user_id': 'user22', 'activity': 'view_item'}
Received: {'user_id': 'user24', 'activity': 'view_item'}
Received: {'user_id': 'user27', 'activity': 'logout'}
Received: {'user_id': 'user28', 'activity': 'logout'}
Received: {'user_id': 'user29', 'activity': 'view_item'}
Received: {'user_id': 'user23', 'activity': 'logout'}
Received: {'user_id': 'user25', 'activity': 'view_item'}
Received: {'user_id': 'user26', 'activity': 'add_to_cart'}
Received: {'user_id': 'user30', 'activity': 'login'}
```

Exercise 3

Create Purchase topic

```
bin/kafka-topics.sh --create \
--topic purchase-topic \
--bootstrap-server localhost:9092 \
--partitions 3 \
--replication-factor 1
```

Create user-activity-topic

```
bin/kafka-topics.sh --create \
--topic user-activity-topic \
--bootstrap-server localhost:9092 \
--partitions 3 \
--replication-factor
```

```
Sent to purchase-topic: {'user_id': 'user1', 'amount': 52.31, 'item': 'backpack'}
Sent to user-activity-topic: {'user_id': 'user9', 'activity': 'login'}
Sent to purchase-topic: {'user_id': 'user6', 'amount': 98.86, 'item': 'laptop'}
Sent to user-activity-topic: {'user_id': 'user3', 'activity': 'login'}
Sent to purchase-topic: {'user_id': 'user1', 'amount': 19.14, 'item': 'headphones'}
Sent to user-activity-topic: {'user_id': 'user2', 'activity': 'purchase'}
Sent to purchase-topic: {'user_id': 'user8', 'amount': 35.16, 'item': 'laptop'}
Sent to user-activity-topic: {'user_id': 'user9', 'activity': 'purchase'}
Sent to purchase-topic: {'user_id': 'user4', 'amount': 81.44, 'item': 'book'}
Sent to user-activity-topic: {'user_id': 'user2', 'activity': 'purchase'}
Sent to purchase-topic: {'user_id': 'user7', 'amount': 38.3, 'item': 'phone'}
Sent to user-activity-topic: {'user_id': 'user1', 'activity': 'login'}
Sent to purchase-topic: {'user_id': 'user3', 'amount': 93.21, 'item': 'book'}
Sent to user-activity-topic: {'user_id': 'user4', 'activity': 'purchase'}
Sent to purchase-topic: {'user_id': 'user8', 'amount': 44.27, 'item': 'backpack'}
Sent to user-activity-topic: {'user_id': 'user1', 'activity': 'login'}
```

With this exercise we evaluated the ability of an producer to feed information into various topics and the advantage of local processment.

```
(venv) joaopino@din-org39 tasks % python3 consumer.py
Listening to user-activity-topic...
```

```
User: user9, Activity Count: 1
User: user3, Activity Count: 1
User: user2, Activity Count: 1
User: user9, Activity Count: 2
User: user2, Activity Count: 2
User: user1, Activity Count: 1
User: user4, Activity Count: 1
User: user1, Activity Count: 2
```

```
(venv) joaopino@din-org39 tasks % python3 consumer.py
Listening to purchase-topic...
User: user1, Total Spent: 52.31
User: user6, Total Spent: 98.86
User: user1, Total Spent: 71.45
User: user8, Total Spent: 35.16
User: user4, Total Spent: 81.44
User: user7, Total Spent: 38.30
User: user3, Total Spent: 93.21
User: user8, Total Spent: 79.43
```

Assignment 2 - Hadoop Map Reduce

For this assignment we had to create a python environment with python 3.10 for it be compatible to MrJob package.

- python3.10 -m venv python3_10_venv
- source python3_10_venv/bin/activate
- pip install mrjob

Exercise 1

- python3 MapReduce/Task_1/task1.py MapReduce/Task_1/fakefriends.csv
- Read the CSV
- Map all entries in key value form:
(age: number of friends)
 - Reducer to Computes the total number of friends and calculates the average number of friends for each age.
 - Converts age to an integer and yields it alongside the average.

Streaming final output from	03	298.2	
57	258.8333333333333	66	276.44444444444446
58	116.54545454545455	67	214.625
59	220.0	68	269.6
60	202.71428571428572	69	235.2
61	256.22222222222223	50	254.6
62	220.76923076923077	51	302.14285714285717
63	384.0	52	340.6363636363636
64	281.3333333333333	53	222.85714285714286
39	169.28571428571428	22	<u>206.42857142857142</u>
40	250.8235294117647	23	246.3
41	268.555555555555554	24	233.8
26	242.05882352941177	25	197.45454545454547
27	228.125	42	303.5
28	209.1	43	230.57142857142858
29	215.916666666666666	44	282.16666666666667
30	235.8181818181818	45	309.53846153846155
31	267.25	46	223.69230769230768
32	207.9090909090909	47	233.22222222222223
33	325.3333333333333	48	281.4
34	245.5	49	184.666666666666666
35	211.625	18	343.375
36	246.6	19	213.27272727272728
37	249.33333333333334	20	165.0
38	193.53333333333333	21	350.875
65	298.2	54	278.0769230769231
		55	295.53846153846155
		56	306.66666666666667

Exercise 2

- python3 MapReduce/Task_2/task2.py MapReduce/Task_2/1800.csv

```
Streaming final output from /var/folders/92/0wqgbrgx6q7b_jn7zqcjctbm000gn/T/task2.joaopino.20241208.211730.882219/output...
"ITIE" -148
"EZE" -135
Removing temp directory /var/folders/92/0wqgbrgx6q7b_jn7zqcjctbm000gn/T/task2.joaopino.20241208.211730.882219...
```

- Mapper: Extracts the capital name and minimum temperature (**TMIN**) as key-value pairs (capital, temperature).
- Reducer: Finds and yields the minimum temperature for each capital.

Exercise 3

- python3 MapReduce/Task_3/task3.py MapReduce/Task_3/Book

- Mapper 1: Extracts words from lines (removing punctuation and converting to lowercase) and emits (word, 1).
- Reducer 1: Counts occurrences of each word and emits (word, total_count).
- Mapper 2: Prepares words for sorting by emitting (None, (count, word)).
- Reducer 2: Sorts words by count in descending order and emits (word, count).

```
"this" 15
"not" 15
"mr" 15
"room" 15
"you" 16
"on" 16
"sir" 19
"so" 19
"but" 22
"for" 22
"up" 22
"mrs" 23
"him" 24
"as" 26
"i" 30
"had" 38
"with" 38
"hall" 39
"it" 42
"at" 44
"her" 45
"that" 48
"said" 52
"in" 57
"was" 71
"of" 78
"she" 79
"his" 84
"he" 89
"to" 91
"a" 108
"and" 161
"the" 230
```

Exercise 4

- python3 MapReduce/Task_4/task4.py MapReduce/Task_4/customer-orders.csv
- Mapper 1: Extracts customer IDs and their spending amounts, emitting (customer_id, amount).
 - Reducer 1: Sums up the spending for each customer, emitting (customer_id, total_amount).
 - Mapper 2: Prepares data for sorting by emitting (None, (total_amount, customer_id)).
 - Reducer 2: Sorts customers by total spending in descending order and emits (customer_id, total_amount).

"68"	6375.449999999997
"73"	6206.199999999999
"39"	6193.109999999999
"54"	6065.389999999999
"71"	5995.660000000003
"2"	5994.59
"97"	5977.189999999995
"46"	5963.109999999999
"42"	5696.840000000003
"59"	5642.89
"41"	5637.62
"0"	5524.949999999998
"9"	5517.240000000001
"85"	5503.43
"61"	5497.479999999998
"32"	5496.050000000004
"58"	5437.730000000005
"63"	5415.150000000001
"15"	5413.510000000001
"6"	5397.879999999998
"92"	5379.280000000002
"43"	5368.83
"70"	5368.249999999999
"72"	5337.44
"34"	5330.8
"9"	5322.649999999999
"55"	5298.090000000002
"90"	5290.409999999998
"64"	5288.689999999996
"93"	5265.750000000001

Assignment 3 - Spark RDD

In this assignment we used the python3.13 env we created in Assignment 1 and added PySpark.

- source venv/bin/activate
- pip install pyspark
- cd RDD_Aproach

In this assignment we didn't add the commands run because it was only necessary to run the python script.

In instances where a new RDD is created thanks to transformations we will only referencing the “Transformation”

Exercise 1

- Read file and create RDD
- Create new RDD from the previous by transforming RDD into a map of tuples: (stationID, entryType, temperature)
- Transform RDD to filter only records where entryType is "TMIN".
- Transform RDD to have tuples (stationID, temperature)
- Reduce using the key stationID by calculating the min in the temperature
- Collects the results and prints

```
('ITE00100554', '-10')
('EZE00100082', '-102')
```

Exercise 2 and 3

- Read file and create RDD
- Transform using flatmap to create more entries based on a split (making the new RDD have all the words)
- Transform to filter RDD to only have words using function isalpha
- Transform RDD to have key-value tuples (word, 1)
- Reduce the RDD based on the key word and summing the values (getting the occurrences of the word)
- [Exercise 3] Transform RDD to be sorted by the key occurrences.
- Collects and prints results

rooms: 1 even: 1 sluggish: 1 Whatever: 1 rid: 1 lot: 1 luggage: 2 hope: 1 stones: 1 aunt: 1 Hastings: 1 swindled: 1 empty: 1 Altogether: 1 vaguely: 1 trudged: 1	had: 38 at: 40 that: 42 in: 53 she: 57 was: 66 he: 67 of: 76 his: 83 to: 90 a: 103 and: 140 the: 212
---	--

Exercise 4 and 5

- Read the file and create the RDD
- Transform the RDD to have key-value pair of (customerID, amount)
- Reduce the RDD by the key customerID summing the value amount
- [Exercise 5] Sort the RDD by the amount
- Collect from the RDD and print the output

Customer_81: 5112.709999999999	Customer_85: 55651.15
Customer_84: 4652.939999999999	Customer_8: 5517.240000000001
Customer_3: 4659.63	Customer_0: 5524.949999999998
Customer_93: 5265.750000000001	Customer_41: 5637.62
Customer_89: 4851.479999999999	Customer_59: 5642.89
Customer_45: 3309.38	Customer_42: 5696.840000000003
Customer_24: 5259.920000000003	Customer_46: 5963.109999999999
Customer_96: 3924.230000000001	Customer_97: 5977.189999999995
Customer_67: 4505.79	Customer_2: 5994.59
Customer_63: 5415.150000000001	Customer_71: 5995.660000000003
Customer_94: 4475.569999999999	Customer_54: 6065.389999999999
Customer_32: 5496.050000000004	Customer_39: 6193.109999999999
Customer_38: 4898.460000000002	Customer_73: 6206.199999999999
Customer_28: 5000.709999999998	Customer_68: 6375.449999999997

Exercise 6 and 7

RDD 1: Marvel Graph

- Read the file and create the RDD
- Transform the RDD by using split to get all the mentions of a superhero
- Transform the RDD by filtering empty words “”
- Transform the RDD creating an map of (word,1)
- Reducing the RDD by the value superheroID and summing the occurrences
- [Exercise 6] Transforming the RDD by filtering from the key occurrence (Descending)
- [Exercise 7] Transforming the RDD by filtering from the key occurrence (Ascending)
- Taking the number 1 spot value
- [After the transformations in RDD 2]
- Collecting the RDD and print the output

RDD 2: Marvel Names

- Read the file and create the RDD
- Transforming the RDD by mapping (id, names) and removing the “ characters from the name
- [After all the transformations on the RDD 1...]
- Transforming the RDD by filtering to only display the value from the superheroID obtained through the collection in RDD1
- Collecting the RDD and print the output

CAPTAIN AMERICA with 1937 occurences

RED WOLF II with 1 occurences

Assignment 4 - Spark SQL

This assignment was very similar to assignment 3 but instead of using RDD we used SQL queries. We created our environment using the following steps:

- source venv/bin/activate
- cd Spark_SQL

Exercise 1

- Create the Spark session and read the csv file containing the weather data
- Create a function called *mapper* that returns the data of each row split in four in order to have each field separated and then proceed with the mapping
- Create a Spark Dataframe after the mapping
- Use spark SQL to create a query that selects the minimum temperature for each station
- Show on the terminal and then stop Spark

```
min_temp_by_station = spark.sql("""  
    SELECT Station, MIN(Temperature) as MinTemperature  
    FROM weather_data  
    GROUP BY Station  
    ORDER BY MinTemperature  
""")
```

Station	MinTemperature
ITE00100554	-148.0
EZE00100082	-135.0
GM000010962	0.0

Exercise 2 and 3

- Create the Spark session and read the file containing the words data
- Create a function called *mapper* that returns the data of each row, in an array, split word by word in lowercase
- Create a Spark Dataframe of the words in the data book after the mapping
- Use spark SQL to create a query that counts the frequency of each word
- In exercise 2 show the frequency of one hundred words without taking account the most frequent words
- In exercise 3 do the same process but this time printing the most frequent words
- Stop Spark

```

query = """
    SELECT word, frequency
    FROM word_frequency
    ORDER BY frequency DESC
"""

```

word	frequency
the	228
and	150
a	106
to	90
he	85
his	84
she	78
of	77
was	67
in	56
at	43
that	42
her	39
had	38
with	38
said	36
it	29
as	25

```

word_frequency = spark.sql("""
    SELECT word, COUNT(*) as frequency
    FROM words_data
    GROUP BY word
""")

```

word	frequency
online	1
biting	1
crest	1
few	2
deftly	1
still	5
eggs,	1
some	5
...	3
filled	1
taking	1
laughing	1
undo.	1
lent	1
approached	1
muffler,	1
afraid,	1

Exercise 4 and 5

- Create the Spark session and read the file containing the consumer data
- Create a function called *mapper* that returns the data of each row in three in order to have each field separated and then proceed with the mapping
- Create a Spark Dataframe of the consumers information retrieved in the last step
- Use spark SQL to create a query that counts the amount spent by each consumer
- In exercise 4 show the amount spent of each consumer without taking account the users that spent the most
- In exercise 5 do the same process but this time printing the top spenders
- Stop Spark

```

total_spent_query = """
SELECT
    customer_ID,
    ROUND(SUM(price), 2) as total_amount_spent
FROM transactions
GROUP BY customer_ID
"""

```

```

total_spent_query = """
SELECT
    customer_ID,
    ROUND(SUM(price), 2) as total_amount_spent
FROM transactions
GROUP BY customer_ID
ORDER BY total_amount_spent DESC
"""

```

customer_ID	total_amount_spent
51	4975.22
7	4755.07
15	5413.51
54	6065.39
11	5152.29
29	5032.53
69	5123.01
42	5696.84
73	6206.2
87	5206.4
64	5288.69
3	4659.63
30	4990.72
34	5330.8
59	5642.89
8	5517.24
22	5019.45
28	5000.71
85	5503.43
35	5155.42

only showing top 20 rows

customer_ID	total_amount_spent
68	6375.45
73	6206.2
39	6193.11
54	6065.39
71	5995.66
2	5994.59
97	5977.19
46	5963.11
42	5696.84
59	5642.89
41	5637.62
0	5524.95
8	5517.24
85	5503.43
61	5497.48
32	5496.05
58	5437.73
63	5415.15
15	5413.51
6	5397.88

only showing top 20 rows

Exercise 6 and 7

- Create the Spark session and read the file containing the data from Marvel superheroes appearances
- Create a variable that counts the appearances of each superhero ID
- Read the file containing the correspondence of each ID to each superhero name
- Create a Spark Dataframe containing the ID and the number of appearances of that ID
- Create a Spark Dataframe of the correspondence of each ID and superhero name
- Use spark SQL to create a query that uses both dataframes in order to obtain the correspondence between the number of appearances and the name of the superhero
- In exercise 6 show the top 20 most popular superheroes
- In exercise 7 do the same process but this time showing the less 20 popular superheroes
- Stop Spark

```

query = """
SELECT n.superhero_name, c.appearances
FROM superhero_names n
JOIN appearance_counts c
ON n.superhero_ID = c.superhero_ID
ORDER BY c.appearances DESC
LIMIT 20
"""

```

superhero_name	appearances
CAPTAIN AMERICA	1933
SPIDER-MAN/PETER PARKER	1741
IRON MAN/TONY STARK	1528
THING/BENJAMIN J. GRANT	1426
WOLVERINE/LOGAN	1394
MR. FANTASTIC/REED RICHARDSON	1386
HUMAN TORCH/JOHNNY SIEGEL	1371
SCARLET WITCH/WANDA MAXIMOFF	1345
THOR/DR. DONALD BLAKEMORE	1289
BEAST/HENRY & HANK PETERSON	1280
VISION	1263
INVISIBLE WOMAN/SUE SPARTACUS	1244
HAWK	1176
CYCLOPS/SCOTT SUMMER	1104
STORM/ORORO MUNROE	1095
ANGEL/WARREN KENNETH MCKEE	1094
WASP/JANET VAN DYNE	1093
ANT-MAN/DR. HENRY JONES	1092
SHE-HULK/JENNIFER WALTERS	1080
DR. STRANGE/STEPHEN STRANGE	1079

```

query = """
SELECT n.superhero_name, c.appearances
FROM superhero_names n
JOIN appearance_counts c
ON n.superhero_ID = c.superhero_ID
ORDER BY c.appearances ASC
LIMIT 20
"""

```

superhero_name	appearances
VINDIKTOR/	1
MARLOPOLIS, EDGAR	1
MURRAY, WILLIAM T. BROWN	1
HOUSTON, LT. COMMANDER	1
SCOPE	1
DESADIA	1
SANDSTORM/	1
GOLEM II	1
BLOWTorch/	1
MASTER OF VENGEANCE	1
PEACEMONGER/	1
PAST MASTER/PROFESSOR	1
SNOW QUEEN/GITTE	1
SPARROW BEAR, MELLIS	1
GAMBIT DOPPELGANGER	1
KURLYCHEK, PATTY	1
GREASE	1
FIREFLY II	1
RIPPER/DR. JACQUELINE	1
MARKS, DR. SHIELA	1